ARTEMIS Foundational Research

Data Collection Findings and Synthesis

FAST NUCES Final-Year Project **F26-008**

Adaptive Reasoning, Task Execution and Multi-agent Intelligence System

Team member	Roll number
Hashir Rauf	23L-2572
Hira Khalid	23L-2594
Doureesha Batool	23L-2651

Supervisor: Mr. Mubasher Hussain

Compiled September 6, 2026

Contents

List of Figures

1 Executive Summary

This report analyses the foundational research collected for ARTEMIS: a questionnaire with 374 valid responses and 13 coded in-person interviews. The purpose is to test whether the problem ARTEMIS targets is real, whether the bounded-scope design is a genuine trust advantage, and which features to build first.

The data confirms the problem. A large majority of respondents work across multiple files often, lose meaningful time reorienting after a break, and see their file organisation decay faster than they clean it. The reasons the work is not already automated are lack of trust in automation and lack of know-how, not indifference. The bounded-workspace idea is supported: roughly two thirds of respondents rate an assistant that only sees files they explicitly add as meaningfully more trustworthy than one with full computer access, and local-first is the single most common deployment preference. Predicted sustained use is moderate rather than enthusiastic, and the most common churn risks are poor fit with existing habits and a mistake the user cannot trust the tool to avoid again.

Headline findings

The nine numbered findings (F1 to F9) are stated in context in the sections that follow. In summary:

- **F1.** The reorientation problem is real. 44% of respondents struggle to remember what they were doing when they return to a project after a break, and only 7% never do.
- **F2.** File disorder outpaces manual cleanup for 59% of respondents (agree or strongly agree).
- **F3.** The problem costs real time at a moderate scale: 42% lose more than an hour a week.
- **F4.** The work is not automated because of distrust of automation (36%) and not knowing how (28%), not because the tasks do not exist.
- **F5.** The bounded-workspace hypothesis holds. 65% rate a workspace-only assistant as more trustworthy than a full-access one (mean 3.8 on a 1 to 5 scale), and 40% prefer local-first over cloud.
- **F6.** Users are forgiving of a single mistake. 65% would keep using ARTEMIS after one wrong suggestion; only 10% would abandon it outright.
- **F7.** The main retention risks are habit fit (27%) and an untrusted mistake (24%); approval fatigue is real but smaller (10%).
- **F8.** Feature demand is concentrated. Session resume, multi-document synthesis, and tidying a folder are the top three first-week features; speech input and output is last by a wide margin.
- **F9.** The interviews confirm the questionnaire and add effortless onboarding as a named condition for continued use.

2 Research Method

2.1 Instruments

The questionnaire had 17 items covering four areas: the respondent's context (role, how often they do multi-file work, prior use of AI assistants); the problem (reorientation after a break, file disorder, time lost, why repeated tasks are not automated); design and trust (approval-first versus act-then-undo, the value of a bounded workspace, local versus cloud, tolerance for a mistake, predicted sustained use, predicted reason to stop); and feature demand (a forced top-three selection from ten candidate features, plus two free-text items). Likert items used a 1 to 5 scale.

The interview guide coded each conversation on five axes: the last time the participant lost their place in multi-file work; the workaround they already use; their reaction to the ARTEMIS concept; what a trust-breaking mistake would do to their usage; and whether they would keep using the tool.

2.2 Sample

Table 1: Response counts

Instrument	Responses
Questionnaire (valid rows)	374
In-person interviews (coded)	13

Recruitment was through FAST NUCES student and practitioner networks. The questionnaire ran from 22 to 27 August 2026. The sample is student-heavy (81% students, 13% developers or software engineers, the remainder freelance or employed), which matches the intended early-adopter profile for a final-year project but limits generalisation.

2.2.1 Limitations

- The sample is self-selected and drawn from a single institution's networks.
- Students dominate; the working-professional signal rests on a smaller subsample.
- ARTEMIS was described in text, not demonstrated. Predicted use and predicted churn are stated intentions, not observed behaviour.
- Free-text responses include a set of near-identical phrasings, which suggests some answer options were presented as suggested text. Theme counts should be read as indicative, not precise.

3 Respondent Profile

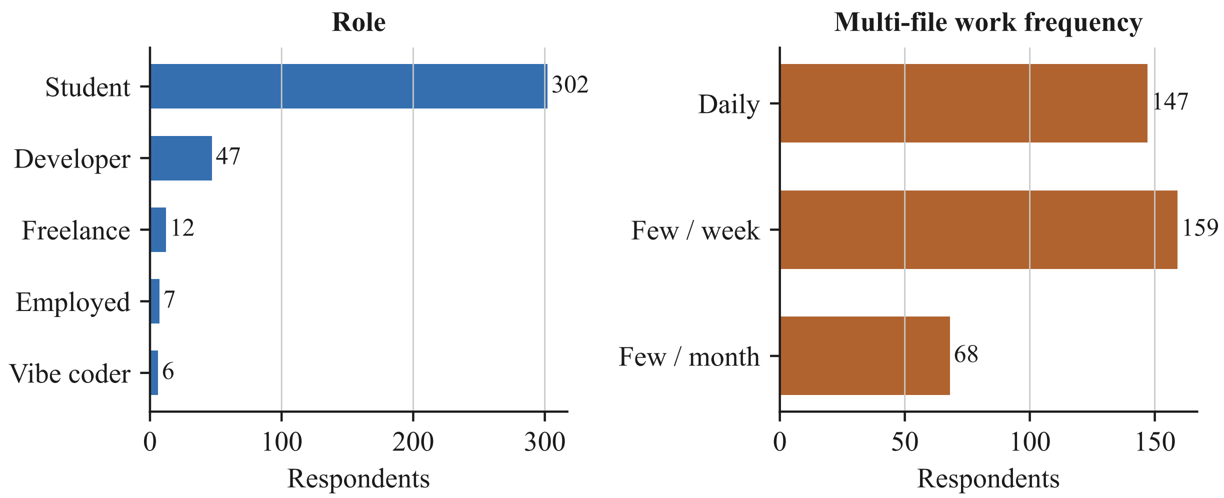

Figure 1: Respondent role and how often respondents do multi-file project work.

3.1 Role and exposure to the problem domain

As shown in Figure 1, 81% of respondents are students and 13% are developers or software engineers. Multi-file work is frequent: 82% do it daily or a few times a week (147 daily), and none of the respondents do it less than a few times a month. The population that answered is therefore squarely inside ARTEMIS's target use case.

3.2 Prior experience with AI assistants

77% have used an AI assistant to help manage files, write code, or handle tasks (regularly or occasionally), and only 10% have never tried one. The audience is experienced enough to compare ARTEMIS against existing tools, which raises the bar for differentiation (Section 6.3).

4 Problem Validation

4.1 Reorientation after a break

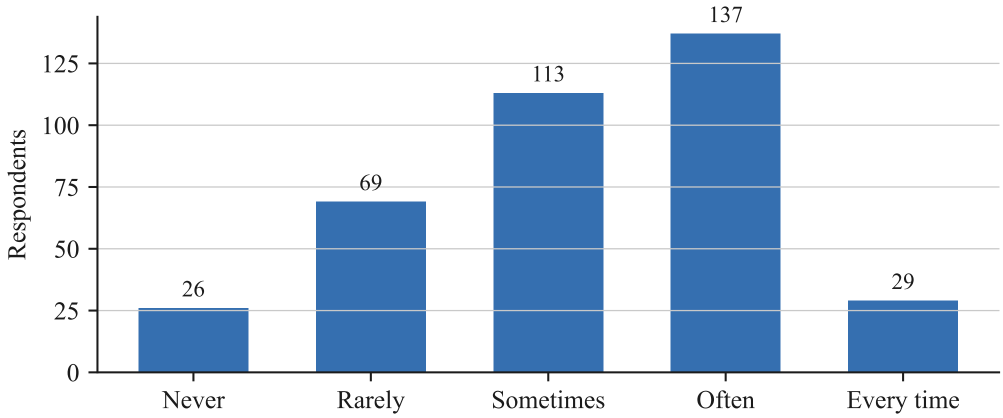

Figure 2: How often respondents return to a project after a break of a day or more and struggle to remember what they were doing or what to do next.

44% of respondents answered "Often" or "Every time" (Figure 2). Only 7% said it never happens. This is the problem the first ARTEMIS feature, session resume, addresses directly.

> **Finding F1.** *The reorientation problem is real and widespread.*
> Nearly half of respondents struggle to recover context after a break, and the overwhelming majority experience it at least sometimes. Session resume targets a real, frequent pain point.

4.2 File disorder outpaces manual cleanup

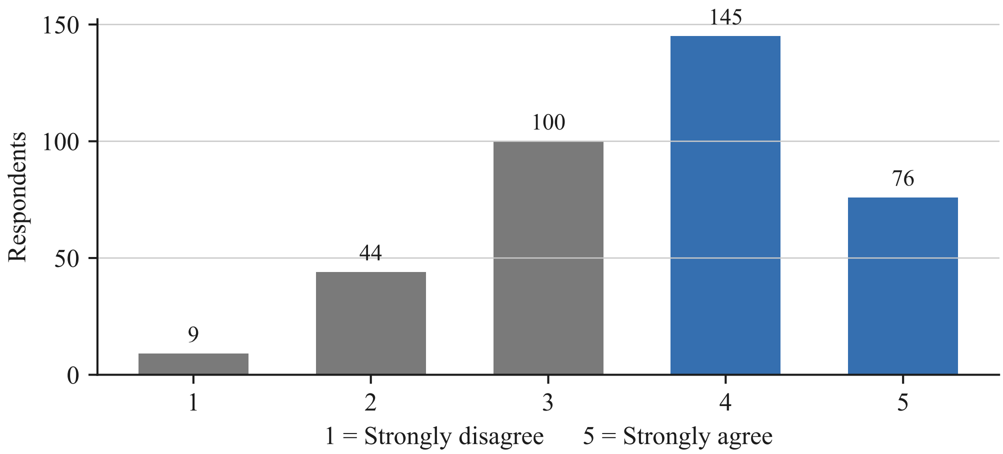

Figure 3: Agreement with "My folders and files tend to become messy or disorganized faster than I can manually clean them up" (1 = strongly disagree, 5 = strongly agree).

59% of respondents rate this 4 or 5 (Figure 3), mean 3.6. The tidy-a-folder feature has a clear constituency, though Section 8 notes a countervailing segment that will not delegate organisation.

> **Finding F2.** *Most users cannot keep up with their own file organisation.*
> A clear majority agree their files decay faster than they can clean them, which validates the tidy-a-folder feature as a genuine need rather than a nice-to-have.

4.3 Self-reported time lost

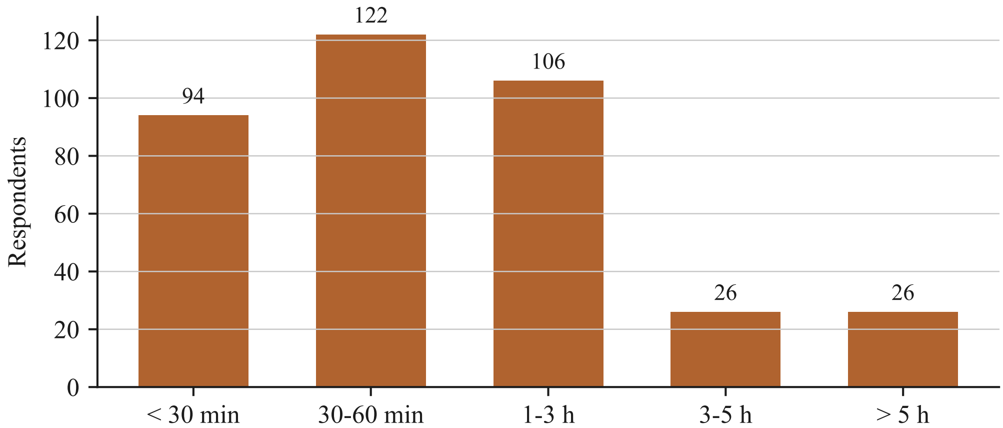

Figure 4: Estimated total time lost per week to reorienting after breaks, tidying files, repeating manual tasks, and searching scattered notes.

42% of respondents estimate losing more than an hour a week, and 14% estimate more than three hours (Figure 4). The modal answer is 30 to 60 minutes. This is a moderate, recurring cost rather than a crisis, which matches the moderate predicted-use numbers in Section 5.

> **Finding F3.** *The problem costs real time, at a moderate scale.*
> Most respondents lose 30 minutes to 3 hours a week. ARTEMIS should be positioned as removing steady friction, not as solving an emergency.

4.4 Why the work is not already automated

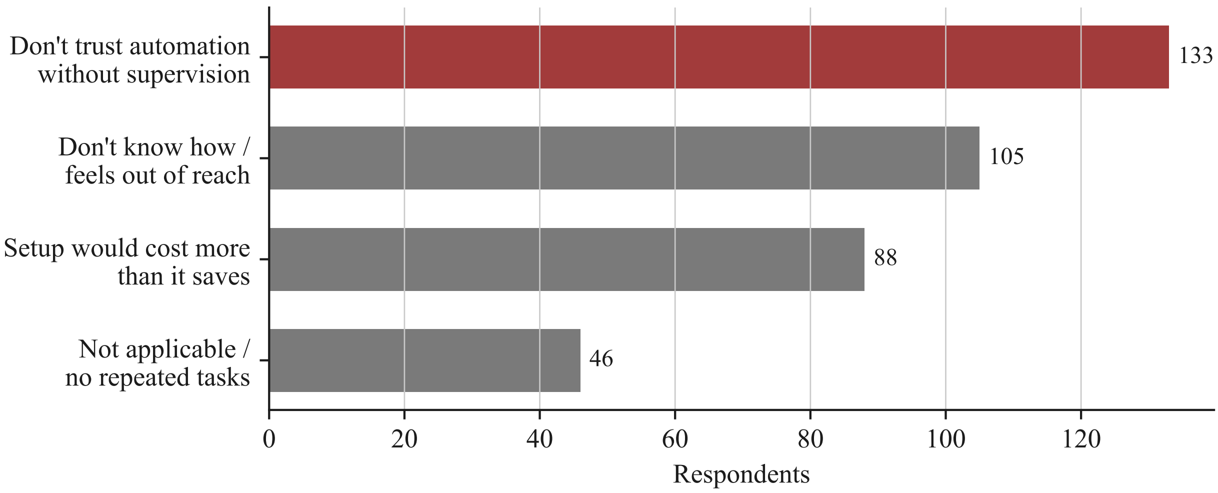

Figure 5: The honest reason respondents have not automated repeated manual file-handling tasks.

The leading reasons are distrust of unsupervised automation (36%) and not knowing how (28%), followed by setup cost exceeding the payoff (24%). Only a minority say the tasks do not exist (Figure 5).

> **Finding F4.** *The barrier is trust and know-how, not absence of the need.*
> Respondents want the work done but do not trust automation to do it unsupervised and do not know how to set it up. An approval-gated assistant with zero configuration directly answers both objections.

5 Design and Trust Signals

5.1 Approval-first versus act-then-undo

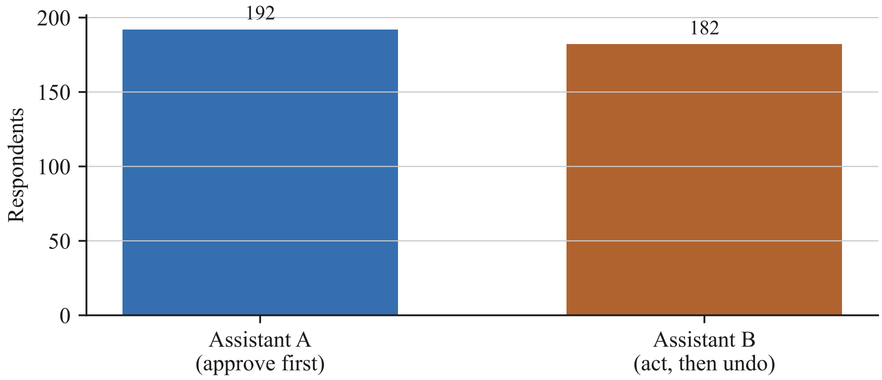

Figure 6: Preference between Assistant A (asks for approval before every file change or outbound action) and Assistant B (acts on reasonable requests automatically, with review and undo afterward).

The split is close: 192 for approval-first, 182 for act-then-undo (51% versus the remainder, Figure 6).

5.1.1 Implication

There is no single right answer for the whole audience. Approval granularity should be a user-configurable setting (for example: approve every action, approve by category, or trust the session after the first few approvals), rather than a fixed product decision. This also mitigates the approval-fatigue churn risk in Section 5.6.

5.2 The bounded-workspace hypothesis

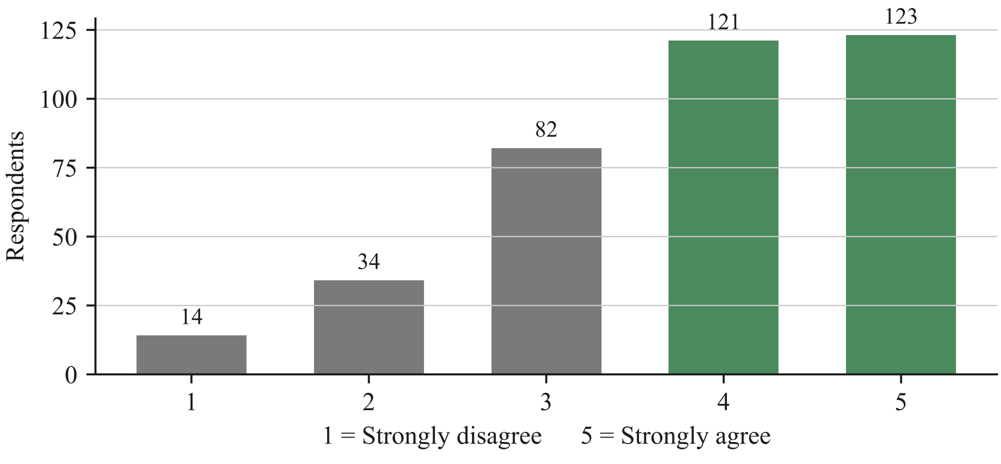

Figure 7: Agreement with "An assistant that only sees files I explicitly add to a workspace is meaningfully more trustworthy to me than one with full access to my computer" (1 = strongly disagree, 5 = strongly agree).

65% of respondents rate this 4 or 5, mean 3.8 (Figure 7). This is the most direct test in the survey of the project's central hypothesis, that restricting an assistant's scope increases user trust, and the result supports it.

Finding F5. *The bounded workspace is a genuine trust advantage.*

Roughly two thirds of respondents find a workspace-only assistant meaningfully more trustworthy than a full-access one. The bounded-scope design should be kept and made visible in the product, not treated as an internal implementation detail.

5.3 Local-first versus cloud

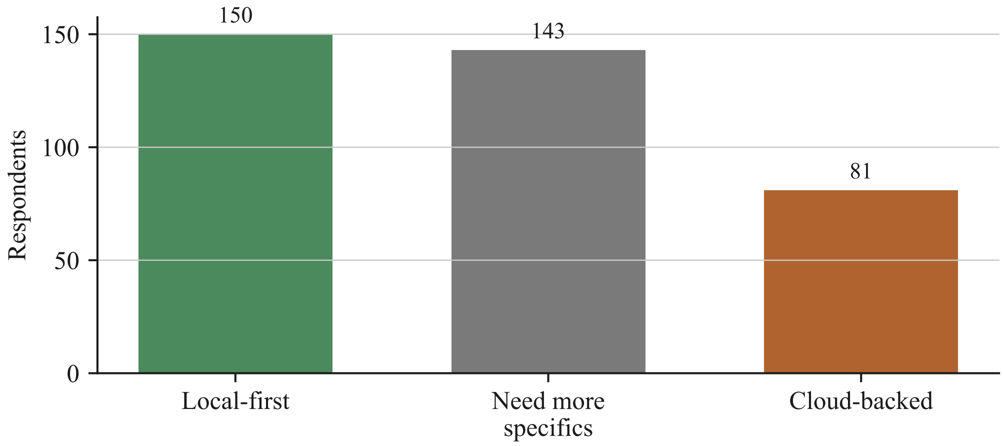

Figure 8: Preference for a local, offline-first version versus a cloud-backed version.

Local-first is the single most common choice at 40%, 38% want more specifics before deciding, and only 22% actively prefer cloud (Figure 8). Privacy also appears as a churn reason (17%, Section 5.6) and as a recurring free-text concern (Section 6.3). Respondents who prefer local-first also report slightly higher predicted sustained use than those who prefer cloud, so the local-first option is not costing engagement.

5.4 Tolerance for a single mistake

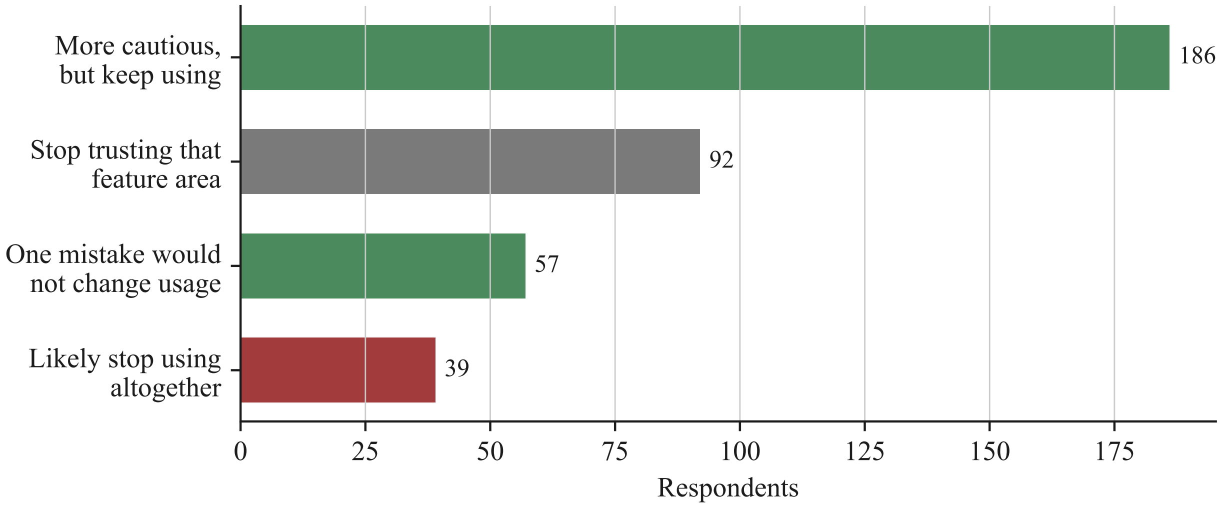

Figure 9: What respondents would do if ARTEMIS made one wrong or unwanted suggestion (for example grouped files incorrectly or drafted an email with a mistake).

65% would keep using ARTEMIS after one mistake (either "more cautious but keep using" or "one mistake would not change my usage"). A further group would stop trusting only the specific feature that erred, and just 10% would abandon the tool entirely (Figure 9).

Finding F6. *One mistake is survivable if recovery is easy.*

Most users are forgiving of a single error, provided they can see and reverse it. This makes a reliable, visible undo more important than eliminating every mistake.

5.5 Predicted sustained use

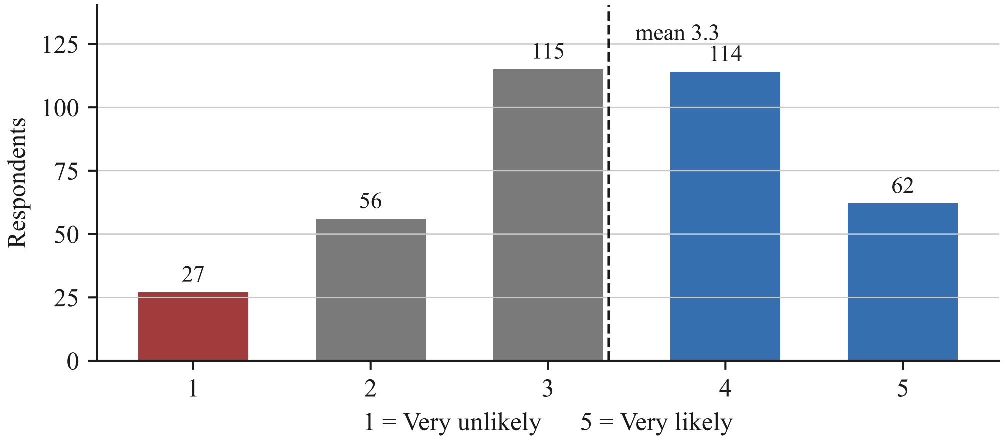

Figure 10: Likelihood of continuing to use ARTEMIS after the novelty wears off, if it worked exactly as described (1 = very unlikely, 5 = very likely).

47% of respondents answered 4 or 5, 22% answered 1 or 2, and the mean is 3.3 (Figure 10). This is a moderate, plausible result for a described-but-not-demonstrated concept, and it sets a realistic expectation: ARTEMIS should aim to convert the large "3" group through a good first experience.

5.6 Predicted reasons for churn

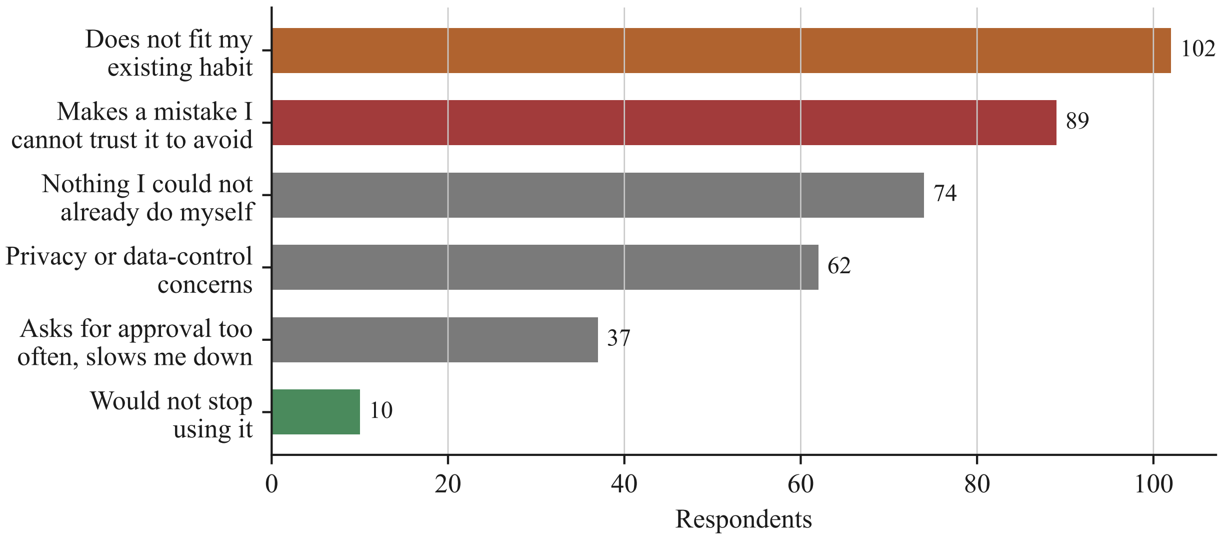

Figure 11: The single most likely reason a respondent would stop using ARTEMIS after a week.

The dominant risks are poor fit with existing habits (27%) and a mistake the user cannot trust the tool to avoid again (24%), followed by redundancy with what the user can already do (20%) and privacy (17%). Approval fatigue is real but the smallest named risk at 10% (Figure 11).

Finding F7. *Retention depends on habit fit and trustworthy recovery, not on fewer prompts.*

The two largest churn drivers are workflow fit and an untrusted mistake. Onboarding that plants ARTEMIS inside an existing routine, plus visible undo, address the top risks. Reducing approval prompts helps a smaller group and should not come at the cost of the trust the approval model buys.

6 Feature Demand

6.1 First-week feature priorities

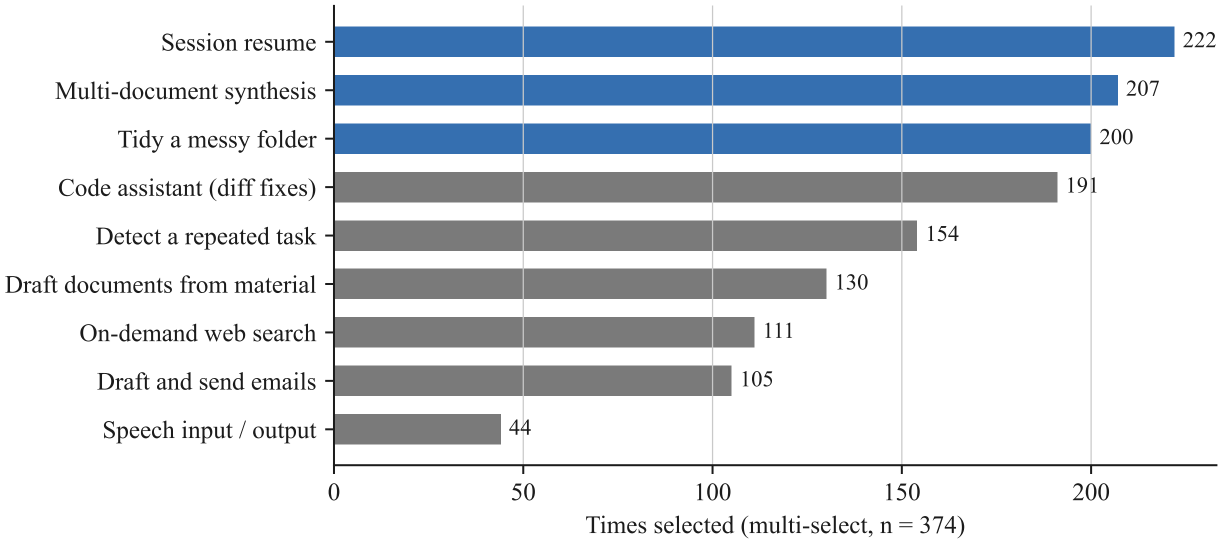

Figure 12: Number of times each feature was chosen in the forced top-three selection (multi-select, 374 respondents).

Ranked by selections (Figure 12):

1. Resuming interrupted work (session resume), 222
2. Pulling together notes and citations from multiple documents, 207
3. Tidying up a messy folder, 200
4. Code assistant (explain errors, suggest fixes as a diff), 191
5. Detecting and offering to automate a repeated task, 154
6. Drafting documents from existing material, 130
7. On-demand web search, 111
8. Drafting and sending emails with approval, 105
9. Speech input and output, 44

Finding F8. *Demand is concentrated on the first three planned features; speech is last.*

Session resume, multi-document synthesis, and folder tidy are the clear top three. Speech input and output is chosen roughly a fifth as often as the leader and should stay last in the build order.

6.2 Alignment with the planned build order

Table 2: Planned ARTEMIS build order against measured first-week demand.

Planned feature (build order)	Demand rank	Comment
Pick up where you left off (session resume)	1	Highest demand and lowest technical risk. Correct to build first.
Tidy up a folder	3	Strong demand; watch the "I organise my own files" segment.
Pull together notes	2	Second-highest demand. Bring forward if resources allow.
Notice a repeated task	5	Mid demand; pairs well with the automation-distrust finding (F4).
Code assistant	4	Solid demand, concentrated among the developer subsample.
Web search	7	Lower demand; keep on-demand only as planned.
Mailing	8	Lower demand; approval and logging are essential given trust concerns.
Document drafting	6	Mid demand.
Speech input and output	9	Lowest demand. Keep last.

The measured demand order is close to the planned build order. The one adjustment worth considering is prioritising multi-document synthesis, which ranks second in demand but is currently scheduled after the file features.

6.3 Unprompted requests

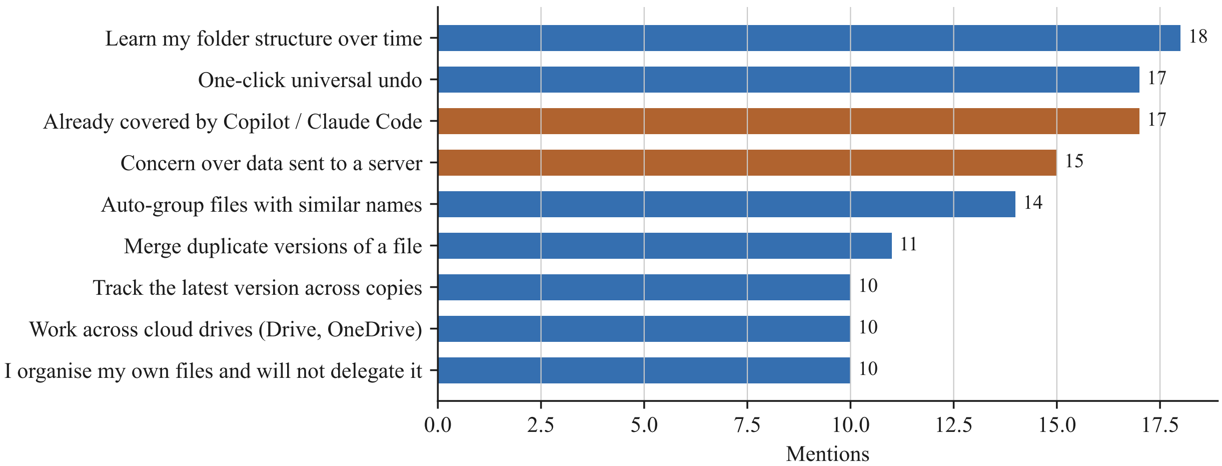

Figure 13: Recurring themes in free-text responses. Counts are indicative given repeated phrasings in the raw data.

The recurring themes, by frequency (Figure 13):

1. **Learn my folder structure over time** (18). The most requested capability not already on the roadmap. Users want ARTEMIS to remember how they organise work so they do not reconfigure it each session.
2. **One-click universal undo** (17). Named repeatedly as the thing that would make the tool trustworthy. Reinforces F5.
3. **Already covered by Copilot or Claude Code** (17). A differentiation risk: an experienced audience is not sure what ARTEMIS adds. The answer is the bounded workspace and local-first operation, and this

needs to be stated explicitly in all messaging.

4. **Concern about how much data is sent to a server** (15). Supports the local-first default and a visible "What ARTEMIS Knows" panel.
5. **File-consolidation requests** (auto-group files with similar names, merge duplicate versions, track the latest version across copies). Together these are the largest cluster and point to a "consolidate related files" capability within the tidy feature.
6. **Work across cloud drives** (Google Drive, OneDrive), not only local folders. In tension with the local-first and bounded-scope design; note for later.
7. **I organise my own files and will not delegate it**. A real segment that will not use the tidy feature regardless of quality.

7 Interview Findings

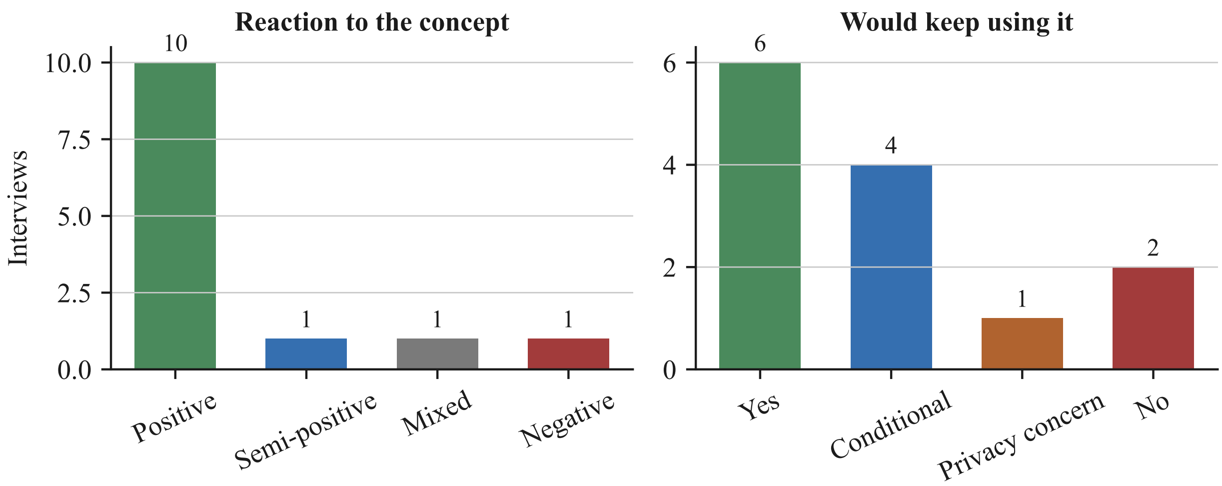

Figure 14: Coded interview responses: reaction to the ARTEMIS concept, and whether the participant would actually keep using it (13 interviews).

7.1 The problem in participants' own words

Almost every participant recognised the problem. The workarounds they already use are consistent: ad hoc Markdown notes written by hand or with an AI agent to record where they left off; manual, repeated restructuring because "every project has a different structure"; and simply tolerating the friction because automating it feels like more work than it saves. Two participants said the problem does not affect them much, one because an IDE agent already handles it.

7.2 Reaction to the concept

10 of 13 participants reacted positively to the concept (Figure 14), with the rest split between conditional, mixed, and one negative.

7.3 Trust-breaking moments

Most participants (8 of 13) said a single mistake would not stop them, though several qualified this with "with caution". A minority wanted the tool to learn from being corrected. 2 said a trust-breaking mistake would make them stop. This mirrors the questionnaire result F5.

7.4 Willingness to keep using

6 said yes outright and 4 said they would keep using it conditionally, with easy onboarding named as the most common condition, followed by privacy and by the tool reaching sufficient maturity.

Finding F9. *Interviews confirm the questionnaire and add onboarding as a named condition.*

Concept reception is positive, a single mistake is mostly survivable, and continued use is conditional on onboarding being effortless.

8 Synthesis

8.1 What the data confirms

- The reorientation and file-disorder problems are real, frequent, and cost most users 30 minutes to 3 hours a week (F1, F2, F3).
- The work is not automated because of distrust and lack of know-how, which an approval-gated, zero-configuration assistant is well placed to address (F4).
- The bounded workspace is a real trust advantage, and local-first is the most common deployment preference (F5).
- Users are forgiving of a single mistake when recovery is easy (F6), and the interviews agree (F9).
- Feature demand is concentrated on session resume, multi-document synthesis, and folder tidy; speech is last (F8).

8.2 What the data challenges

- The audience is split almost evenly on approval-first versus act-then-undo, so a single fixed interaction model will not satisfy everyone.
- A visible segment states plainly that it organises its own files and will not delegate that, which caps the addressable audience for the tidy feature.
- An experienced audience is not sure ARTEMIS adds anything over Copilot or Claude Code. Differentiation is not self-evident and must be argued.
- Predicted sustained use is moderate (mean 3.3), not enthusiastic. The large neutral group is the one to win.

8.3 Cross-cutting tensions

- **Approval friction versus trust.** The approval model is what makes the tool trustworthy, but approval fatigue is a named churn risk. Resolve with configurable granularity, not by weakening the model.
- **Local-first versus cross-drive convenience.** The privacy-driven local-first preference conflicts with requests to operate across Google Drive and OneDrive.
- **Personalisation versus bounded scope.** The top unmet request, learning the user's folder structure over time, implies persistent memory of the user's habits, which must be designed to stay inside the bounded-scope guarantee.

8.3.1 Segment view

Table 3: Selected measures by respondent segment.

Segment	*n*	Workspace-trust (mean, 1–5)	Prefer approval-first	Prefer local-first
All respondents	374	3.8	51%	40%
Students	302	3.8	51%	41%
Developers / engineers	47	3.7	51%	36%
Regular AI-assistant users	172	3.9	44%	40%
Light or non AI-assistant users	87	3.8	59%	41%

Students and developers give near-identical workspace-trust ratings and identical approval-first shares, so role does not predict the interaction preference. Experience does move it: regular AI-assistant users prefer approval-first less often than light or non-users (44% versus 59%), which is consistent with heavier users being more comfortable delegating. The effect is modest and neither group is close to unanimous, so the practical conclusion is unchanged: the interaction model should be a setting rather than a fixed decision. Local-first preference is broadly consistent across all segments (36% to 41%).

9 Recommendations for the Build

1. **Build session resume, multi-document synthesis, and folder tidy first; keep speech input and output last.** Motivated by F8. Consider moving multi-document synthesis ahead of the later file features given it ranks second in demand.
2. **Ship a reliable, visible, one-click undo before any outward-action feature.** Motivated by F6 and the free-text requests. Undo is the single most cited condition for trusting the tool.
3. **Make onboarding effortless and design it to plant ARTEMIS inside an existing routine.** Motivated by F7 and F9, which name onboarding as the top condition for continued use.
4. **Keep the bounded opt-in workspace and offer a local-first mode, and make both visible in the product.** Motivated by F5 and the privacy churn and free-text signals. This is also the differentiation answer to the "already covered by Copilot" objection.
5. **Make approval granularity user-configurable** (per action, per category, or session trust after the first approvals). Motivated by the near-even A versus B split and the approval-fatigue churn risk.
6. **Add "learn my workspace structure over time" to the roadmap.** It is the top requested capability not currently planned. Design it as per-workspace memory of the user's organisation preferences so it stays within the bounded-scope guarantee.
7. **State the differentiation from Copilot and Claude Code explicitly** in the proposal, the UI, and any external material. The target audience already uses those tools and will not infer the difference.
8. **Treat file consolidation as a first-class part of the tidy feature** (group files with similar names, flag duplicate versions, identify the latest version), since together these are the largest free-text cluster.

9.1 Open questions carried forward

- How to expose configurable approval granularity without making the first-run experience complex.
- Whether and how to support cloud drives without breaking the local-first and bounded-scope guarantees.
- How persistent "learn my structure" memory is represented in the "What ARTEMIS Knows" panel.
- How to convert the large neutral group (predicted-use rating 3) through the first session. These extend the existing internal open-questions note.

A Appendix

A.1 Questionnaire items

1. Timestamp.
2. Email address.
3. How often do you work on projects that involve managing multiple files, documents, or folders (coursework, code projects, client deliverables, research)?
4. What best describes you?
5. How often do you return to a project after a break of a day or more and struggle to remember what you were doing or what to do next?
6. "My folders and files tend to become messy or disorganized faster than I can manually clean them up" (1 = strongly disagree, 5 = strongly agree).
7. If you repeat the same manual file-handling task often, why have you not automated it? Pick the honest reason.
8. Roughly how much total time per week do you think you lose to these issues (reorienting after breaks, tidying files, repeating manual tasks, searching scattered notes)?
9. Have you used an AI assistant (Copilot, ChatGPT, Gemini, a computer-use agent) to help manage files, write code, or handle tasks on your computer?
10. Which of these two assistants would you actually prefer to use day to day? (A: asks for approval before every file change or outbound action; B: acts on reasonable requests automatically, with review and undo afterward.)
11. "An assistant that only sees files I explicitly add to a workspace is meaningfully more trustworthy to me than one with full access to my computer" (1 = strongly disagree, 5 = strongly agree).
12. Suppose ARTEMIS made a wrong or unwanted suggestion once (grouped your files incorrectly, or drafted an email with a mistake). Would you keep using it?
13. Which of the following features would you actually use in your first week, if it worked exactly as described? Select your top 3.
14. A local, offline-first version versus a cloud-backed version. Which would you actually choose?
15. If ARTEMIS worked exactly as described, how likely would you be to actually use it for your own work, not "try once" but keep using after the novelty wears off? (1 = very unlikely, 5 = very likely).
16. What is the single most likely reason you would stop using ARTEMIS after a week, if you had to guess right now?
17. Is there anything about how you currently manage multi-file work, a frustration, a workaround, or a feature you wish existed, that this questionnaire did not ask about but that we should know before building a tool like this?

A.2 Interview coding scheme

Table 4: Interview coding axes.

Axis	Coded values
The last time you lost your place	Problem exists / happened a lot / never happened / other free-text descriptions.
The workaround you already have	Markdown notes, manual restructuring, "too lazy to automate", not equipped, no need.
Reacting to the concept	Positive / semi-positive / mixed / negative.
The trust-breaking moment	Shrug it off (sometimes with caution) / conditional / needs human feedback / stop using.
Would you actually keep using it	Yes / conditional (onboarding, maturity) / privacy concern / no.

A.3 Complete response distributions

Counts and percentages for every closed questionnaire item and every coded interview axis.

Table 5: Q4. What best describes you?

Response	Count	Percent
Student	302	80.7%
Developer / Software Engineer	47	12.6%
Freelance	12	3.2%
Employed	7	1.9%
Vibe coder	6	1.6%
Total	374	100.0%

Table 6: Q3. Frequency of multi-file project work

Response	Count	Percent
Daily	147	39.3%
A Few Times a Week	159	42.5%
A Few Times a Month	68	18.2%
Total	374	100.0%

Table 7: Q9. Prior use of an AI assistant for files / code / tasks

Response	Count	Percent
Yes, regularly	172	46.0%
Yes, occasionally	115	30.7%
Tried it once or twice	48	12.8%
Never	39	10.4%
Total	374	100.0%

Table 8: Q5. Return after a break and struggle to remember what to do next

Response	Count	Percent
Never	26	7.0%
Rarely	69	18.4%
Sometimes	113	30.2%
Often	137	36.6%
Every time	29	7.8%
Total	374	100.0%

Table 9: Q6. Files become messy faster than I can clean them up (1 to 5)

Response	Count	Percent
1	9	2.4%
2	44	11.8%
3	100	26.7%
4	145	38.8%
5	76	20.3%
Total	374	100.0%

Table 10: Q8. Estimated time lost per week to these issues

Response	Count	Percent
Under 30 minutes	94	25.1%
30 minutes - 1hour	122	32.6%
1-3 hours	106	28.3%
3-5 hours	26	7.0%
Over 5 Hours	26	7.0%
Total	374	100.0%

Table 11: Q7. Why repeated manual file tasks are not automated

Response	Count	Percent
I don't trust automation to do it correctly without supervision	133	35.8%
I don't know how / it feels out of reach	105	28.2%
It would take longer to set up than the time it saves	88	23.7%
Not applicable ; I don't have repeated manual tasks	46	12.4%
Total	372	100.0%

Table 12: Q10. Preferred assistant: approve-first (A) vs act-then-undo (B)

Response	Count	Percent
Assistant A (approve first)	192	51.3%
Assistant B (act, then undo)	182	48.7%
Total	374	100.0%

Table 13: Q11. A workspace-only assistant is more trustworthy than a full-access one (1 to 5)

Response	Count	Percent
1	14	3.7%
2	34	9.1%
3	82	21.9%
4	121	32.4%
5	123	32.9%
Total	374	100.0%

Table 14: Q14. Local-first vs cloud-backed preference

Response	Count	Percent
Local First even if its slower or weaker	150	40.1%
I would need to know more specifics before deciding either way	143	38.2%
Cloud-backed even if its less private	81	21.7%
Total	374	100.0%

Table 15: Q12. Reaction to one wrong or unwanted suggestion

Response	Count	Percent
I'd be more cautious but keep using it	186	49.7%
I'd stop trusting its suggestions in that specific feature area	92	24.6%
Yes, one mistake wouldn't change my usage	57	15.2%
I'd likely stop using the tool altogether	39	10.4%
Total	374	100.0%

Table 16: Q15. Likelihood of sustained use if it worked as described (1 to 5)

Response	Count	Percent
1	27	7.2%
2	56	15.0%
3	115	30.7%
4	114	30.5%
5	62	16.6%
Total	374	100.0%

Table 17: Q16. Single most likely reason to stop using ARTEMIS after a week

Response	Count	Percent
I forget to use it/ it doesn't fit into my existing habit	102	27.3%
It makes a mistake I dont trust it to avoid again	89	23.8%
It doesn't do anything I couldn't already do myself easily	74	19.8%
Privacy or data-control concerns	62	16.6%
It asks for approval too often and slows me down	37	9.9%
I don't think I would stop using it	10	2.7%
Total	374	100.0%

Table 18: Q13. Features that would actually be opened in the first week (pick top 3)

Response	Count	Percent
Resuming interrupted work (session resume)	222	16.3%
Pulling together notes/citations from multiple documents	207	15.2%
Tidying up a messy folder	200	14.7%
Code assistant (explain errors, suggest fixes as a diff)	191	14.0%
Detecting and offering to automate a repeated task	154	11.3%
Drafting documents from existing material	130	9.5%
On-demand web search	111	8.1%
Drafting and sending emails (with approval)	105	7.7%
Speech input/output	44	3.2%
Total	1364	100.0%

Table 19: Q17. Recurring themes in open-ended feedback

Response	Count	Percent
Learn my folder structure over time	18	14.8%
One-click universal undo	17	13.9%
Already covered by Copilot / Claude Code	17	13.9%
Concern over data sent to a server	15	12.3%
Auto-group files with similar names	14	11.5%
Merge duplicate versions of a file	11	9.0%
Track the latest version across copies	10	8.2%
Work across cloud drives (Drive, OneDrive)	10	8.2%
I organise my own files and will not delegate it	10	8.2%
Total	122	100.0%

Table 20: Interview: reaction to the concept

Response	Count	Percent
Positive	10	76.9%
Semi-positive	1	7.7%
Mixed	1	7.7%
Negative	1	7.7%
Total	13	100.0%

Table 21: Interview: would you actually keep using it

Response	Count	Percent
Yes	6	46.2%
Conditional	4	30.8%
Privacy concern	1	7.7%
No	2	15.4%
Total	13	100.0%

Table 22: Interview: reaction to a trust-breaking moment

Response	Count	Percent
Shrug it off	8	61.5%
Conditional	2	15.4%
Stop using	2	15.4%
Needs human feedback	1	7.7%
Total	13	100.0%